


- ▷ Yogesh Simmhan
- ▷ simmhan@iisc.ac.in
- ▷ Department of Computational and Data Sciences
- ▷ Indian Institute of Science, Bangalore



DS256 (3:1)

Scalable Systems for Data Science




Distributed Systems

Helping Data Science Scale

Vertical and Horizontal Scaling

- ▷ *How do you get better performance for an application?*
 - Better algorithm
 - Better programming
 - Better hardware
- ▷ **Vertical Scaling (Scale up)**
 - Adding more hardware resources to an existing machine
 - Faster CPU, more cores, more memory
 - Simpler. Little to no (multi-threaded or loosely coupled) application constraint
 - Upper bound of hardware limits
- ▷ **Horizontal Scaling (Scale Out)**
 - Adding more number of machines
 - Cumulatively more number of CPU cores, memory
 - Applications need to be designed to use multiple machines

Scale Up (*Faster CPUs/Clockspeed*)




Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
 New plot and data collected for 2010-2019 by K. Rupp

Scale Up (More CPU cores)

AMD EPYC Zen

▷ Multi/Many-core processing

- Multiple cores on single machine
 - Shared local memory space
 - Cores may be heterogeneous
- Independent execution on each core
 - A *single thread/process does not run faster*
- Multiple processes on a desktop/laptop
- Multi-threaded applications on a server




<https://www.nextplatform.com/2017/06/20/competition-returns-x86-servers-epyc-fashion/>

<https://community.arm.com/developer/ip-products/processors/b/processors-ip-blog/posts/arm-dynamIQ-technology-for-the-next-era-of-compute>

<https://www.tomshardware.com/news/annotated-apple-m3-processor-die-shots-bring-chip-designs-to-life>

<https://developer.nvidia.com/blog/inside-nvidia-blackwell-ultra-the-chip-powering-the-ai-factory-era/>


<https://www.nextplatform.com/2024/10/10/amd-turns-the-screws-with-turin-server-cpus/>



AMZ Epyc Zen5 Turin


Scale Up (SoC Heterogeneous cores)

- ▷ Has a flavor of scale out
 - Applications may need to be redesigned to exploit parallelism
 - May be able to increase to 100s of cores



Apple M3

ARM big.LITTLE



<https://www.nextplatform.com/2017/06/20/competition-returns-x86-servers-epyc-fashion/>
<https://community.arm.com/developer/ip-products/processors/b/processors-ip-blog/posts/arm-dynamIQ-technology-for-the-next-era-of-compute>
<https://www.tomshardware.com/news/annotated-apple-m3-processor-die-shots-bring-chip-designs-to-life>
<https://developer.nvidia.com/blog/inside-nvidia-blackwell-ultra-the-chip-powering-the-ai-factory-era/>
<https://www.nextplatform.com/2024/10/10/amd-turns-the-screws-with-turin-server-cpus/>

Scale Out

- ▷ More processing units, working *collaboratively*
- ▷ **Distributed Systems**
 - Multiple servers, connected over the network
 - **Distributed Computing**
 - Clusters of machines, with CPU/GPU/Memory
 - Distributed memory space
 - **Distributed Storage**
 - Disks attached to clusters of machines
 - Network Attached Storage
- ▷ *How can we make effective use of multiple machines?*



Distributed Systems

- ▷ **Commodity clusters & Clouds**
 - Available off the shelf at large volumes, lower cost
 - Virtualized resources, modest performance
 - “Low speed” Ethernet: Higher latency, lower bandwidth
 - Loosely coupled applications



Param Pravega cluster@SERC
Ground Floor

- ▷ **HPC clusters and Supercomputers**
 - High-end CPUs
 - Fast network interconnect. Low latency, high bandwidth.
 - Tightly coupled applications

How many cores? PF?

30k Cores
3.5PF

- ▷ *How can we make effective use of many machines of modest capability?*

Cloud Computing

- ▷ Large Data Centers
 - 1000's of racks, 100k servers
- ▷ Virtualized infrastructure
 - On-demand: Rent VMs by the minute
 - 100 machines for 10mins costs same as 1 machine for 1000 minutes
- ▷ Makes use of economies of scale.
Much cheaper than on-premises servers.
 - Reduce operations costs (energy, personnel)
 - Ensures capital costs (servers) are fully used



Ashburn, Virginia

Designing applications for distributed systems

- ▷ Software is critical to enable use of distributed computing and storage
 - System software for resource management
 - Data management systems
 - Distributed programming models for designing apps
 - Runtime environments to execute these apps
- ▷ Question: How do you design a (C++/Java/Python) program to add a billion numbers, from scratch?

Scalability and Parallel Speedup

Jun and Ni, JPDC 1993

- ▷ **Speedup:** Performance measure for parallel processing
 - Ratio of execution time on a *uniprocessor* and on a *parallel processor*
- ▷ Different variations of Speedup, depending on the *assumptions* and *constraints* on application and system
 - Fixed problem size
 - Fixed time for processing
 - Fixed amount of memory
 - Effect of uneven workload
 - Effect of communication overhead



Scalable Multi Processor

- ▷ Scalable Multi Processor
 - As **number of processors** increase, the **amount of memory** and the ***network bandwidth*** also increases
 - Homogeneous processors
 - Shared or distributed memory
 - Different network topologies
 - Message passing or share memory programming paradigms

Speedup on N processors

- ▷ W is amount of **work** for application
- ▷ $T_i(W)$ is **time** to complete W amount of **work** on i **processors**
- ▷ Speedup on N processors with W amount of work is:

$$S_N(W) = \frac{T_1(W)}{T_N(W)}$$

Degree of Parallelism for an Application

- ▷ Maximum number of processors that can be actively processing at a given phase of the application
- ▷ Amount of parallel processing that can be exploited by the application during its execution lifetime
 - Load can be uneven over its lifetime
- ▷ Parallelism Profile and Shape
 - *For each degree of parallelism, what's the duration?*




FIG. 1. Parallelism profile of an application.





FIG. 2. Shape of the application.

Uneven Workloads: Different degrees of parallelism for different phases



Parallel Speedup: **One** processor with **uneven** load, no communication

- ▷ W_i is amount of work available with a degree of parallelism i
 - Across all tasks of the application
- ▷ m is maximum degree of parallelism for application
- ▷ Total work for application is:

$$W = \sum_{i=1 \text{ to } m} W_i$$

- ▷ Δ is computing capacity of one processor
- ▷ Execution time of W_i on a **single processor** is $t_1(W_i) = \frac{W_i}{\Delta}$
- ▷ **Total execution time** of W on a single processor is:

$$T_1(W) = \sum_{i=1 \text{ to } m} t_1(W_i) = \sum_{i=1 \text{ to } m} \frac{W_i}{\Delta}$$

Parallel Speedup: **Infinite** processors with **uneven** load, no communication

- ▷ With i processors, the execution time for W_i is $t_i(W_i) = \frac{W_i}{i \cdot \Delta}$
- ▷ With ∞ processors, the execution time for W_i still remains:

$$t_{\infty}(W_i) = \frac{W_i}{i \cdot \Delta}, 1 \leq i \leq m$$

- *Parallelism benefits accrue only up to the degree of parallelism that is available for a task*

- ▷ Total execution time of W on ∞ processors is:


$$T_{\infty}(W) = \sum_{i=1 \text{ to } m} t_{\infty}(W_i) = \sum_{i=1 \text{ to } m} \frac{W_i}{i \cdot \Delta}$$

Parallel Speedup: Infinite processors with **uneven** load, no communication

- ▶ Maximum speedup with work W and ∞ processors:

$$S_{\infty}(W) = \frac{T_1(W)}{T_{\infty}(W)} = \frac{\sum_{i=1 \text{ to } m} \frac{W_i}{\Delta}}{\sum_{i=1 \text{ to } m} \frac{W_i}{i \cdot \Delta}} = \frac{\sum_{i=1 \text{ to } m} W_i}{\sum_{i=1 \text{ to } m} \frac{W_i}{i}}$$

- ▶ S_{∞} is the **best possible speedup** based on inherent parallelism of the application
 - Machine-dependent factors are ignored



Parallel Speedup: **N** processors with **uneven** load, no communication

- With N processors and i degree of parallelism, the load on each processor will be different if $N < i$

- $W_i \cdot \left\lfloor \frac{i}{N} \right\rfloor$, for less loaded ones
- $W_i \cdot \left\lceil \frac{i}{N} \right\rceil$, for more loaded ones

$$i = 7, N = 3 \\ 2/7 \text{ vs. } 3/7$$



- Time elapsed for W_i will be limited by the **slowest** of these
- With N processors, the execution time for W_i is:

$$t_N(W_i) = \frac{W_i}{i \cdot \Delta} \cdot \left\lceil \frac{i}{N} \right\rceil$$


- Total execution time of W on N processors is:

$$T_N(W) = \sum_{i=1 \text{ to } m} \frac{W_i}{i \cdot \Delta} \cdot \left\lceil \frac{i}{N} \right\rceil$$

Parallel Speedup: **N** processors with **uneven** load, no communication

- ▶ Speedup with work W and N processors is:

$$S_N(W) = \frac{T_1(W)}{T_N(W)} = \frac{\sum_{i=1 \text{ to } m} W_i}{\sum_{i=1 \text{ to } m} \frac{W_i}{i} \cdot \left\lceil \frac{i}{N} \right\rceil}$$



Mapping “tasks” to **N=3** processors

Next task phase can start only after previous phase has completed

Note: This model consolidates phases with the same degree of parallelism, i.e., W_i across all phases is consolidated and run in parallel. So its an approximation.

Parallel Speedup: **N** processors with **uneven** load, **with** communication

- ▷ Communication latency is machine dependent, and application dependent
 - NW topology, routing, etc.
- ▷ If the communication cost for N processors to complete W work is $Q_N(W)$
- ▷ Assuming degree of parallelism is not affected by communication costs
- ▷ Speedup with work W and N processors is:

$$S_N(W) = \frac{T_1(W)}{T_N(W)} = \frac{\sum_{i=1 \text{ to } m} W_i}{\left(\sum_{i=1 \text{ to } m} \frac{W_i}{i} \cdot \left\lceil \frac{i}{N} \right\rceil \right) + Q_N(W)}$$

Parallel Speedup with Constraints: Fixed Problem Size

- ▷ Fixed problem size constraint
 - Size of the application is constant, W
 - Amount of work to be done is constant
 - E.g., number of iterations, size of the matrix, input file size
 - Need to complete the application as quickly as possible
 - *How does the **fixed-work speedup** vary with increasing the number of processors?*
- ▷ This is the default formulation, same as above:

$$S_N(W) = \frac{T_1(W)}{T_N(W)} = \frac{\sum_{i=1 \text{ to } m} W_i}{\left(\sum_{i=1 \text{ to } m} \frac{W_i}{i} \cdot \left[\frac{i}{N} \right] \right) + Q_N(W)}$$

Parallel Speedup with Constraints: **Fixed Time**

- ▷ Fixed time constraint
 - **Time budget** for processing the application is fixed
 - Usually set as the time taken for application of some size to complete on a single processor
 - **Size of the application is increased with the increase in the number of processors**
 - E.g., more iterations, larger matrix size, larger input file
 - Need to complete the scaled application within time budget
 - *How does the **fixed-time speedup** vary with increasing the number of processors?*

Parallel Speedup with Constraints: Fixed Time

- ▷ W' is the “scaled” amount of work for an application, when given N processors
- ▷ m' is the maximum degree of parallelism for the scaled problem
 - This can be different from m
- ▷ Total amount of scaled work is $W' = \sum_{i=1 \text{ to } m'} W'_i$
- ▷ For fixed-time, the constraint imposed is: $T_1(W) = T_N(W')$
 - *Time to run W on 1 processor is same as time to run W' on N processors*
- ▷ This is re-written as:

$$\sum_{i=1 \text{ to } m} W_i = \left(\sum_{i=1 \text{ to } m'} \frac{W'_i}{i} \cdot \left\lceil \frac{i}{N} \right\rceil \right) + Q_N(W')$$

Parallel Speedup with Constraints: Fixed Time

- ▷ Speedup with fixed-time is:

$$S_N(W') = \frac{T_1(W')}{T_N(W')} = \frac{T_1(W')}{T_1(W)} = \frac{\sum_{i=1 \text{ to } m'} W'_i}{\sum_{i=1 \text{ to } m} W_i}$$

- ▷ This is the time to complete the scaled problem on one processor divided by the non-scaled problem on one processor

Simplified Models of Speedup

- ▷ Negligible communication costs, $Q_N = 0$
- ▷ Workload only has a **sequential** and a “perfectly parallel” part, i.e., W_i is non-zero only for:
 - $i = 1$ (sequential)
 - $i = N$ (perfectly parallel)
- ▷ Sequential part is independent of system size
 - Remains constant even when problem is “scaled”
 - $W_1 = W'_1$

Simplified Models of Speedup: Fixed Problem Size




FIG. 3. Amdahl's law.

- ▶ Fixed-size speedup reduces to:

$$S_N(W) = \frac{\sum_{i=1 \text{ to } m} W_i}{\left(\sum_{i=1 \text{ to } m} \frac{W_i}{i} \cdot \left\lceil \frac{i}{N} \right\rceil \right) + Q_N(W)}$$

$$= \frac{W_1 + W_N}{W_1 + \frac{W_N}{N}}$$

- ▶ **Amdahl's Law of Parallel Speedup**
 - For large values of N , the sequential part will dominate the total time taken, $W_1 \gg \frac{W_N}{N}$
 - Maximum Speedup is bound by $\frac{W_1 + W_N}{W_1}$

Amdahl's Law of Strong Scaling


- ▷ Amdahl's Law for Application Scalability
 - Total problem size is fixed
 - *Speedup limited by sequential bottleneck*
- ▷ s is *serial* fraction of application, cannot be parallelized
- ▷ $(1 - s)$ is fraction of application that can be *parallelized*
- ▷ p is number of processors
- ▷ t is time taken to process input of size x on 1 processor
- ▷ Then the speedup is limited by the serial fraction

$$\text{Speedup}(p, x) = \frac{\text{time}(1, x)}{\text{time}(p, x)} = \frac{t}{s.t + \frac{(1-s)}{p}.t} = \frac{s.t + \frac{(1-s)}{1}.t}{s.t + \frac{(1-s)}{p}.t} = \frac{1}{s + \frac{(1-s)}{p}}$$


- ▷ For large values of N , the sequential part will dominate, and maximum speed up is $\frac{1}{s + \frac{(1-s)}{p}} \approx \frac{1}{s}$

s is *serial* fraction of application, cannot be parallelized
 $(1 - s)$ is fraction of application that can be *parallelized*

$$\frac{1}{s + \frac{(1-s)}{p}}$$



Amdahl's Law (Strong Scaling)



© Martha A. Kim

Meaningful only for small values of p , small fractions of s

© Daniels220 at English Wikipedia

Simplified Models of Speedup:

Fixed Time




FIG. 4. Gustafson's scaled speedup.


- ▷ Fixed-time constraint becomes: $W_1 + W_N = W'_1 + \frac{W'_N}{N}$
- ▷ Since sequential part is constant even with scaled problem $W_1 = W'_1$, we have $W_N = \frac{W'_N}{N}$, i.e., $W'_N = N \cdot W_N$
- ▷ Fixed-time speedup reduces to:
$$S_N(W') = \frac{W_1 + N \cdot W_N}{W_1 + W_N}$$
- ▷ **Gustafson's Law of Weak Scaling**

Gustafson's Law of weak scaling


- ▷ Gustafson's Law of Application Scalability
 - Problem size ($p.x$) increases with number of processors (p)
 - “Scaled speedup”
- ▷ $Speedup(p, x.p) = \frac{time(1,p.x)}{time(p,p.x)} = \frac{s.p.t + (1-s).p.t}{s.p.t + \frac{(1-s).p.t}{p}}$
- ▷ If $s.p.t \approx s.t$...i.e., only parallelizable fraction of input work increases with number of processors
- ▷ $Speedup(p, x.p) = \frac{s.t + (1-s).p.t}{s.t + \frac{(1-s).p.t}{p}} = \frac{s + (1-s).p}{s + (1-s)}$
 $= s + (1 - s).p$

Gustafson's Law of weak scaling

- As input work increases, all incremental work can be parallelized. Total time remains constant as we increase number of processors.



Gustafson's Law: $S(P) = P \cdot a \cdot (P-1)$



Scalability

- ▷ **Speedup Efficiency:** How good is the speedup relative to perfect linear speedup?

$$\text{Speedup Efficiency} = \frac{\text{Speedup}(p, _)}$$
$$p$$

- ▷ **Strong Scaling:** How the performance varies with the # of processors for a *fixed total problem size*
- ▷ **Weak Scaling:** How the performance varies with the # of processors for a *fixed problem size per processor*
 - Big Data platforms are intended for “Weak Scaling”

Additional Reading

“

- ▷ Computer Architecture and Amdahl's Law, Gene M. Amdahl, *IEEE Computer*, 2013, doi:10.1109/MC.2013.418
- ▷ Reevaluating Amdahl's Law, Gustafson, John L., *Communications of the ACM*, 31 (5), 1988, doi:10.1145/42411.42415
- ▷ The evolution of distributed computing systems: from fundamental to new frontiers. Lindsay, D., Gill, S.S., Smirnova, D. et al. *Computing* (2021).
<https://doi.org/10.1007/s00607-020-00900-y>
- ▷ Big Data computing and clouds: Trends and future directions, M. D. Assunção, et al., *Journal of Parallel and Distributed Computing*, Volumes 79–80, 2015, Pages 3-15, <https://doi.org/10.1016/j.jpdc.2014.08.003>

“

End of Lecture 2

- ▷ REMINDER
- ▷ *Read the GFS paper*