

- ▷ Yogesh Simmhan
- ▷ simmhan@iisc.ac.in
- ▷ Department of Computational and Data Sciences
- ▷ Indian Institute of Science, Bangalore

DS256 (3:1)

Scalable Systems for Data Science




Module 2

Processing Large Volumes of Big Data

Role of Big Data Processing System for Large Data Volumes

Big Data (Volume)



Required Reading

“

- *MapReduce: Simplified Data Processing on Large Clusters*, Dean and Ghemawat, USENIX OSDI, 2004
- *Resilient distributed datasets: A fault-tolerant abstraction for in-memory cluster computing*, Matei Zaharia, et al., USENIX NSDI, 2012
- *Spark SQL: Relational Data Processing in Spark*, Michael Armbrust, et al., ACM SIGMOD 2015
 - Select chapters from *Learning Spark*, Holden Karau, et al., 1st Editions and *Learning Spark*, Jules S. Damji, Brooke Wenig, Tathagata Das, Denny Lee, 2nd Edition
 - Optional: Clash of the Titans: MapReduce vs. Spark for Large Scale Data Analytics, Juwei Shi, Yunjie Qiu, Umar Farooq Minhas, Limei Jiao Chen Wang, Berthold Reinwald, and Fatma Ozcan, VLDB 2015

MapReduce

MapReduce: Simplified Data Processing on Large Clusters, Dean and Ghemawat, USENIX OSDI, 2004

MapReduce

*“A simple and powerful interface that
enables automatic parallelization and
distribution of large-scale computations,
combined with an implementation of this
interface that achieves high performance on
large clusters of commodity PCs.”*

Google's Cluster Model (2004)

- ▷ Machines have **2 x86 CPU**, running Linux, with **2-4 GB of memory**
 - 2020: Raspberry Pi 4B has 4-core ARM, 2-4GB RAM, Rs.5000 ☺
- ▷ Commodity networking – **100 Mbps or 1 Gbps** per machine, low bisection bandwidth.
- ▷ A cluster has **100's – 1000's** of machines, failures are common
- ▷ Storage on inexpensive IDE disks attached to servers
 - **GFS manages the data**
- ▷ Users submit jobs to a scheduling system
 - Each job is a set of tasks
 - Mapped by the scheduler to available machines in cluster


MapReduce Design Pattern

- ▷ Programming model for distributed applications
 - Clean abstraction for programmers
 - Automatic parallelization & distribution
- ▷ Fault-tolerance
- ▷ Batch data processing system
 - Large inputs sizes
- ▷ Simple data-intensive applications
 - Distributed Grep: Document list → Occurrence of search term
 - URL Access Frequency: URL access list → <URL, freq>
 - Reverse Web-Link Graph: <target,src> → <src, target[]>
 - Term-Vector per Host: <host,word[]> → <host,<word,freq>[]>

MapReduce: Data-parallel Programming Model


- ▷ Process data using **map** & **reduce** user-defined functions
- ▷ **map(k_i , v_i) → List< k_m , v_m >**
 - *map* is called once on every input item
 - Emits a series of intermediate key/value pairs
- ▷ **shuffle & sort phase**
 - All map output values (v_m) with a given key (k_m) are *grouped* together, keys *sorted* within a group
 - Happens internally within the framework
- ▷ **reduce(k_m , List< v_m >) → List< k_r , v_r >**
 - *reduce* is called once on *every unique key & all its values*
 - Emits a value that is added to the output

Inverted Index using MR



- ▷ **map(url, line)**
 - *foreach(word in line.split()) emit(word, url)*
- ▷ **reduce(word, url[])**
 - *save(word, findDistinct(url[]))*

Map-Shuffle-Sort-Reduce



Histogram using MR

7	2	11	2
2	1	11	4
9	10	6	6
6	3	2	8
0	5	1	10
2	4	8	11
5	0	1	0
M M M M			
1,1	0,1	2,1	0,1
0,1	0,1	2,1	1,1
2,1	2,1	1,1	1,1
1,1	0,1	0,1	2,1
0,1	1,1	0,1	2,1
0,1	1,1	2,1	2,1
1,1	0,1	0,1	0,1
Shuffle			
2,1	0,1 0,1	1,1	
2,1	0,1 0,1	1,1	
2,1	0,1 0,1	1,1	
2,1	0,1 0,1	1,1	
2,1	0,1 0,1	1,1	
2,1	0,1 0,1	1,1	
2,1	0,1 0,1	1,1	
2,1	1,1		
R R R			
2,8	0,12	1,8	

```

int bucketWidth = 4 // input

Map(k, v) {
    emit(floor(v/bucketWidth), 1)
    // <bucketID, 1>
}

```

```


// one reduce per bucketID
Reduce(k, v[]){
    sum=0;
    foreach(n in v[])
        sum+=n;
    emit(k, sum)
    // <bucketID, frequency>
}

```

Data transfer &
shuffle between
Map & Reduce
(28 items)


Limitations of MapReduce

- ▷ Multi-stage computing not simple
 - Many different jobs
- ▷ Complex code for simple transformations
 - Repetitive, not *data centric*



Limitations of MapReduce

- ▷ Limited support for non-text, Non-static data




General Batch Processing

Specialized Systems:

iterative, interactive, streaming, graph, etc.

Limitations of MapReduce

- ▷ Poor performance for:
 - Complex, multi--stage applications (e.g. iterative machine learning & graph processing)
 - Interactive *ad hoc* queries



Latency & Bandwidth


- ▷ L1 cache reference
- ▷ L2 cache reference
- ▷ **Main memory reference**
- ▷ **Send 1K bytes over 1 Gbps network**
- ▷ **Read 4K randomly from SSD***
- ▷ **Read 1MB sequentially from memory**
- ▷ **Round trip within same datacenter**
- ▷ **Read 1MB sequentially from SSD***
- ▷ **Send 1MB over 1 Gbps network**
- ▷ **Disk seek**
- ▷ **Read 1MB sequentially from disk**
- ▷ **Send packet CA->NL->CA**

Latency & Bandwidth

▷ L1 cache reference	0.5	ns			
▷ L2 cache reference	7	ns			
▷ Main memory reference	100	ns			
▷ Send 1K bytes over 1 Gbps network	10,000	ns	10	μs	
▷ Read 4K randomly from SSD*	150,000	ns	150	μs	
▷ Read 1MB sequentially from memory	250,000	ns	250	μs	
▷ Round trip within same datacenter	500,000	ns	500	μs	
▷ Read 1MB sequentially from SSD*	1,000,000	ns	1,000	μs	1 ms
▷ Send 1MB over 1 Gbps network	8,250,000	ns	8,250	μs	8 ms
▷ Disk seek	10,000,000	ns	10,000	μs	10 ms
▷ Read 1MB sequentially from disk	20,000,000	ns	20,000	μs	20 ms
▷ Send packet CA->NL->CA	150,000,000	ns	150,000	μs	150 ms

Latency Numbers Every Programmer Should Know


2020





Bandwidth of Memory >> Network or Disk

MTTF in Data Center



"The MTBF (mean time between failures) across all data centers we investigate (with hundreds of thousands of servers) is only 6.8 minutes, while the MTBF in different data centers varies between 32 minutes and 390 minutes."

→ **MTBF with 1000 servers is 680mins**
→ **MTBF with 100 servers is 6800mins (4.7 days)**





Figure 4: Kaplan Meier survival estimate of datacenter switches, shaded region shows 95% confidence intervals.



*Failures may be infrequent during the lifetime
of an application execution*

From MapReduce to Spark

- ▷ Google's MapReduce
 - Programming Model
 - Apache Hadoop runtime environment



Apache Spark

Resilient distributed datasets: A fault-tolerant abstraction for in-memory cluster computing, Zaharia, Chowdhury, Das, Dave, Ma, McCauly, Franklin, Shenker and Stoica, USENIX NSDI, 2012

Learning Spark

Holden Karau, Andy Konwinski, Patrick Wendell & Matei Zaharia,
O'Reilly, First and Second Editions

The Spark Ecosystem

- ▶ Core Spark Engine
 - RDDs, Transformations, Actions, batch processing
- ▶ Higher level abstractions
 - Data frames, SQL-like queries
 - Discretized streams, semi-realtime data
 - Machine learning libraries, **Mllib**
 - Linked data analytics, **GraphX**





Figure 1-1. The Spark stack

Spark: A Distributed Execution Engine


- ▷ **Driver:** User program for application, uses Spark Context, local variables
- ▷ **Spark Context:** Gives access to distributed computing environment
- ▷ **Worker:** Machines on which actual heavy-lift happens
- ▷ **Executor:** Spark execution environment in a worker, Process, exclusive to an application
- ▷ **Task:** Single operation on data, thread



Spark RDD

Resilient Distributed Dataset

- ▶ Collection of homogeneous objects
 - Order is not preserved*
- ▶ Distributed on workers
 - 1 or more Partitions
- ▶ Read-only, immutable
- ▶ Can be rebuilt
- ▶ Can be cached
- ▶ MR like data-parallel operations
 - Execute on workers



Creating and Operating on an RDD

- ▷ Users can provide driver code in multiple languages:
Scala, Java, Python (PySpark)
 - Spark offers equivalent transformations and actions in each language
- ▷ Actual Spark execution environment is in Scala
- ▷ Create RDD by loading data
 - HDFS, local vars, filesystem, NoSQL DB, etc.
 - Data is loaded on partitions on different workers
- ▷ RDD Object offers a logical view of the dataset
 - Can perform operations on the object

Example 3-1. Creating an RDD of strings with `textFile()` in Python

```
>>> lines = sc.textFile("README.md")
```

Creating and Operating on an RDD

- ▷ Users can provide driver code in multiple languages:
Scala, Java, Python (PySpark)
 - Spark offers equivalent transformations and actions in each language
- ▷ Actual Spark execution environment is in Scala
- ▷ Create RDD by loading data
 - HDFS, local vars, filesystem, NoSQL DB, etc.
 - Data is loaded on partitions on different workers
- ▷ RDD Object offers a logical view of the dataset
 - Can perform operations on the object

Example 3-1. Creating an RDD of strings with `textFile()` in Python

```
>>> lines = sc.textFile("README.md")
```

Language Bindings

- ▷ Users can provide driver code in multiple languages
 - Scala, Java, Python
- ▷ Spark offers equivalent transformations and actions in each language
- ▷ Logic within transformations and actions can also be in these languages
- ▷ Actual Spark execution environment is in Scala
 - Standard data structure mapping
 - Python code is pickled (de/serialize) and shipped remotely

Passing Functions to Spark Operations

- ▷ Lambda syntax
 - Functions are *input parameters* to other functions
 - Pass short functions concisely, inline

```
pythonLines = lines.filter(lambda line: "Python" in line)
```



```
def hasPython(line):  
    return "Python" in line
```

```
pythonLines = lines.filter(hasPython)
```

Programming with RDDs

Learning Spark

Holden Karau, Andy Konwinski, Patrick Wendell & Matei Zaharia,
O'Reilly, First Edition

Chapter 3

Basics of Transformations

- ▷ Returns a new RDD, computed lazily
- ▷ Transforms tend to be **element-wise** operations
 - Iterate through each item, apply the operation, e.g. *Filter*
- ▷ **Filter** on *inputRDD* does not affect *inputRDD*
 - Returns a new RDD, *warningsRDD*
- ▷ **Union** operates on two RDDs
 - One of them is an input parameter

```
inputRDD = sc.textFile("log.txt")
errorsRDD = inputRDD.filter(lambda x: "error" in x)
warningsRDD = inputRDD.filter(lambda x: "warning" in x)
badLinesRDD = errorsRDD.union(warningsRDD)
```

Basics of Actions

- ▷ Actually triggers operations, returns a final result to driver
 - Force any required transformations to be executed
 - Count, Take, Collect
- ▷ Result of action must fit in memory of driver
 - Else, can write RDD to HDFS, saveAsTextFile
- ▷ RDDs are computed from scratch when actions are called...See *persist/cache*

```
print "Input had " + badLinesRDD.count() + " concerning lines"
print "Here are 10 examples:"
for line in badLinesRDD.take(10):
    print line
```

Lazy Evaluation

- ▷ Transformations are lazily evaluated
 - Calling a transform does NOT immediately execute it
- ▷ Action triggers execution of *dependent transformations*
- ▷ E.g., `load().map().count()`
 - Load & Map do not execute till we see Count
- ▷ Allows Spark to reduce the number of passes through the data
 - Materializes RDD only when required
 - Reused RDDs that have been materialized earlier
 - Immutability!

Lineage Graph

- ▷ Keeps track of operations used to derive an RDD
- ▷ Helps *lazily materialize* RDD
- ▷ Helps *recover* RDD or their partitions that are lost

