

- ▷ Yogesh Simmhan
- ▷ simmhan@iisc.ac.in
- ▷ Department of Computational and Data Sciences
- ▷ Indian Institute of Science, Bangalore

DS256 (3:1)

Scalable Systems for Data Science



Module 2



Processing Large Volumes of Big Data

Spark Internals

Select topics from external sources

Spark Architecture

- ▷ Spark runs on a cluster of machines
- ▷ Driver interfaces with SparkContext
- ▷ Logical Plan
 - Converts application to a dataflow of dependencies
- ▷ Physical Plan
 - Converts dataflow into specific tasks for execution
 - Tasks executed within Workers/Executors



Logical Plan


Narrow and Wide Dependencies

▷ Narrow Dependency

- Each output partition depends on exactly one or a few input partitions
- Each input partition is used by exactly one output partition
- So output and input partitions can be on same Worker
- Also called full dependency
- E.g., map, cogroup


▷ Wide Dependency

- Each output partition depends on parts of one or more input partition
- Each input partition can be used by one or more output partitions
- So forces a shuffle across Workers
- Also called Shuffle or partial dependency
- E.g., join, groupByKey



Hash Shuffle in Spark


- ▷ Write to one file per reducer per map task



Sort Shuffle in Spark >= v1.2.0


(~MapReduce)

- ▷ Write to a single (spill) file per Map task, maintain offset for each reducer



Narrow Dependency (One to One)


- ▷ 1:1 mapping between output and input records, e.g., map, filter
- ▷ 1:1 mapping between output and input partitions, e.g., mapPartitions



Narrow Dependency (Range)



- ▷ Range dependency
 - Ranges are retained between input and output
 - E.g., union

union(): RangeDependency




Narrow Dependency (N:1)

- ▷ Many to One Dependency
 - Multiple input partitions generate a single output partition
 - Operates on multiple RDDs
 - Requires both RDDS to
 - Have same number of partitions
 - Use same partitioner
 - E.g., cogroup, coalesce
- ▷ If # parts or partitioner not same, this can become wide dependency




Wide Dependency (Single RDD)

- ▶ Each input partition can contribute to multiple output partitions
- ▶ Since data moves from an input to multiple output workers, this causes a shuffle
- ▶ **Shuffles are costly!**




- ▶ Shuffle can be across partitions for one RDD
 - E.g., groupByKey, reduceByKey, sort, distinct




Wide Dependency (Single RDD)



reduceByKey(f, numPartitions)



Example (WordCount): `reduceByKey(_ + _, 2)`




Wide Dependency (Single RDD)




Wide Dependency (Single RDD)

sortByKey(ascending, numPartitions)





Example of sortByKey(true, 2)





Wide Dependency (Multiple RDDs)


- ▷ Shuffle can also combine partitions from multiple RDDs
 - E.g., join, intersect



Wide Dependency (Multiple RDDs)




Physical Plan and Scheduling



Logical Plan



- ▷ Convert logical plan to execution on cluster
- ▷ Create a **Job** for each **Action**
- ▷ Use DAG scheduler to create stages, **Tasks** for Job
- ▷ Schedule and coordinates tasks on cluster




Jobs: Stages and Tasks

DAG Scheduler

- ▶ Stages are separated by a wide dependency
- ▶ Creates tasks for different stages
 - ShuffleStage
 - ResultStage
- ▶ Each Task is responsible for one output partition
- ▶ Task set is a collection of tasks to generate an output RDD




ComplexJob
including map(), partitionBy(), union(), and join()





- DAG Scheduler **schedules** tasks onto Executors
- Triggers subsequent stage once tasks of previous stage complete

- Narrow dependencies are pipelined within a single Task
 - Reduces scheduling overheads, data copies, barrier sync delays




Worker and Executor Model

- ▶ **Workers host one or more Executors**
 - Worker is one machine
 - Executor is one Process
- ▶ **Executors execute Tasks**
 - Each task is responsible for one output partition
- ▶ **Tasks run on separate Threads**
 - Thread pool for concurrency
 - Low thread overhead for executing a task




Execution Model

- ▶ Executors manage **blocks**, with DAG Scheduler
- ▶ Executors manage block **persistence**
- ▶ Executors coordinate shuffle




Executors Coordinate Shuffle



Shuffle

- ▷ Allows items from one partition to be used by multiple partitions
- ▷ Shuffle sits at a stage boundary
- ▷ Requires coordination between tasks sets at the boundary of two stages
- ▷ Shuffle is costly
 - Requires a lot of disk I/O
 - Requires a lot of Network communication
 - Requires a barrier synchronization

Shuffle write in Worker Node (2 cores, 4 ShuffleMapTasks, 3 reducers, consolidateFiles = true)



Spark Tuning

Spark Tuning

- ▶ Improving Parallelism
 - Data parallel execution reduces overall time
- ▶ More partitions allows data parallel execution
 - Number of tasks map to number of partitions
 - Number of threads map to number of tasks
- ▶ Control the number of partitions,
 - numPart params for wide transforms




Figure 7-3. Relationship of Spark tasks, cores, partitions, and parallelism

Data Partitioning

- ▷ RDDs are spread on partitions across nodes of a cluster
- ▷ Controlling the **number of partitions**
 - Helps balance compute load
 - Helps reduce *shuffle* communication across nodes
- ▷ **Partitioning function** to map key to specific partitions
 - Set of “related” keys are placed in same partition
 - E.g., Hash partitioning, Range partitioning
- ▷ Can also be used for **specific algorithms**
 - Operate on partition at a time



Tracking Partitioners

- ▷ Spark **tracks partitioner** used to generate an RDD
 - Uses this to optimize operations
- ▷ A partitioner option **may or may not be present**
 - *cogroup, joins, group/reduce/combineByKey, partitionBy, sort* will set a partitioner
 - Get the partitioner used to generate an RDD using **partitioner()** method
- ▷ Use partitioner to **optimize** future operations
 - E.g., Join output has been hash partitioned. So joined RDD tagged with that partitioner. Helps future **reduceByKey**.
- ▷ Many transformations can **unset** the partitioner
 - Map following a Join can unset Hash partitioner
 - Filter, mapValue & flatMapValues retain partitioner

Impact of Partitioning: Limit Shuffle

- ▶ Join of big RDD with small RDD
 - Default: All to all shuffle without any partitioner
 - HashPartition: Selectively move small RDD to big RDD

```
val userData = sc.sequenceFile[UserID, UserInfo]("hdfs://...").persist()  
val events = sc.sequenceFile[UserID, LinkInfo](logFileName)  
val joined = userData.join(events)  
  
val userData = sc.sequenceFile[UserID, UserInfo]("hdfs://...")  
    .partitionBy(new HashPartitioner(100)) // Create 100 partitions  
    .persist()
```




Passing Partitioners

- ▷ Can explicitly call **partitionBy** transformation
- ▷ Pass a number of partitions/partitioner as part of wide transformations
 - E.g., *join*, *groupByKey*, *sort*, etc.
- ▷ Default Partitioners
 - **HashPartitioner**: Uses key/value's hash to decide partition
 - **RangePartitioner**: Creates roughly equal ranges, determined by sampling the RDD contents
- ▷ Custom Partitioners
 - **numPartitions**: number of partitions to create
 - **getPartition(key)**: Returns 0 to (numPartitions-1)

Tuning Spark

- ▷ Static vs. Dynamic resource allocation
 - Static fixes number of executors at submission time
 - Guarantees fixed resources before job starts
 - Dynamic allows resources to scale out and in by up to min/max executors
 - Based on task queue demand, idle time threshold
- ▷ Spark Memory Allocation
 - Execution memory used by shuffle, join
 - Storage memory used by user data, to cache blocks/partitions
 - Reserved memory is buffer to avoid out of memory exceptions



Spark Tuning

- ▷ Caching and Persistence
- ▷ Caching attempts to retain partitions in memory
 - Use when RDDs used across actions
 - Limited by available memory across executors
 - Recreated if OOM on executors
- ▷ Persist on disk allows moving to disk
 - “Unlimited” space, but slower than cache
- ▷ Cache/persist for iterative operations on RDDs, common RDDs used by several transformations
- ▷ Do no cache if too big for memory, or cheap to recompute

Additional Reading

- ▷ **Spark Internals**
 - Lijie Xu (Jerry Lead)
 - <https://github.com/JerryLead/SparkInternals>
- ▷ **The Internals of Apache Spark 3.1.1**
 - Jacek Laskowski
 - <https://books.japila.pl/apache-spark-internals/overview/>
- ▷ **Learning Spark**, Holden Karau, Andy Konwinski, Patrick Wendell & Matei Zaharia, O'Reilly
 - Chapter 7 (2nd Ed)