

- ▷ Yogesh Simmhan
- ▷ simmhan@iisc.ac.in
- ▷ Department of Computational and Data Sciences
- ▷ Indian Institute of Science, Bangalore

DS256 (3:1)

Scalable Systems for Data Science



Module 2

Processing Large Volumes of Big Data

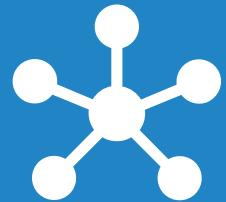


“

- ▷ *Programming Assignment 1 on Spark Data Frames posted on 7 Feb.*
 - ▷ *Due on 28 Feb.*
 - ▷ *Teams of two*
 - ▷ *20% weightage*
- ▷ *No external help, GenAI coding, etc.*

“

- ▷ Quiz 1 on Modules 1 & 2 on Thu, 12 Feb
 - ▷ All topics till today's class



Spark DataFrames & SQL

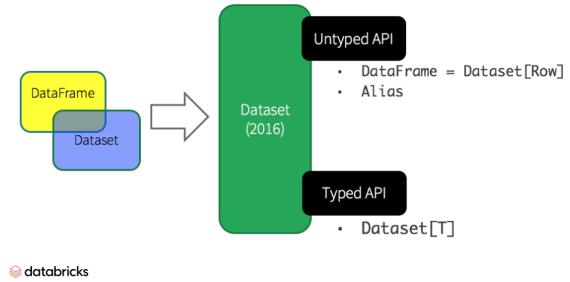
Spark SQL: Relational Data Processing in Spark,
Michael Armbrust, et al., ACM *SIGMOD* 2015

Limitations of Spark RDD

- ▷ Spark only offers high-level constructs on iterations over RDD items, invocation pattern
- ▷ Lambda expressions or functions are opaque to Spark
 - No optimizations of execution
- ▷ Types are opaque (other than being homogeneous)
 - No type-specific behavior
- ▷ Limitations on efficient execution
 - **Imperative programming:** Users tell how to execute
 - Relies on users to optimize code

DataFrames

- ▷ DataFrame: Inspired by [pandas](#)
 - RDD: Inspired by Python native operators like map
- ▷ More expressive and simpler
 - Compose a SQL-like query...
 - Using high-level DSL operators and APIs
- ▷ Tell Spark *what* to do
 - Spark can parse query, understand our intention
 - Optimize/arrange operations for efficient execution
- ▷ Even more uniformity across language bindings
 - Avoid *user-defined code!*
- ▷ Allows you to drop down to RDD also!

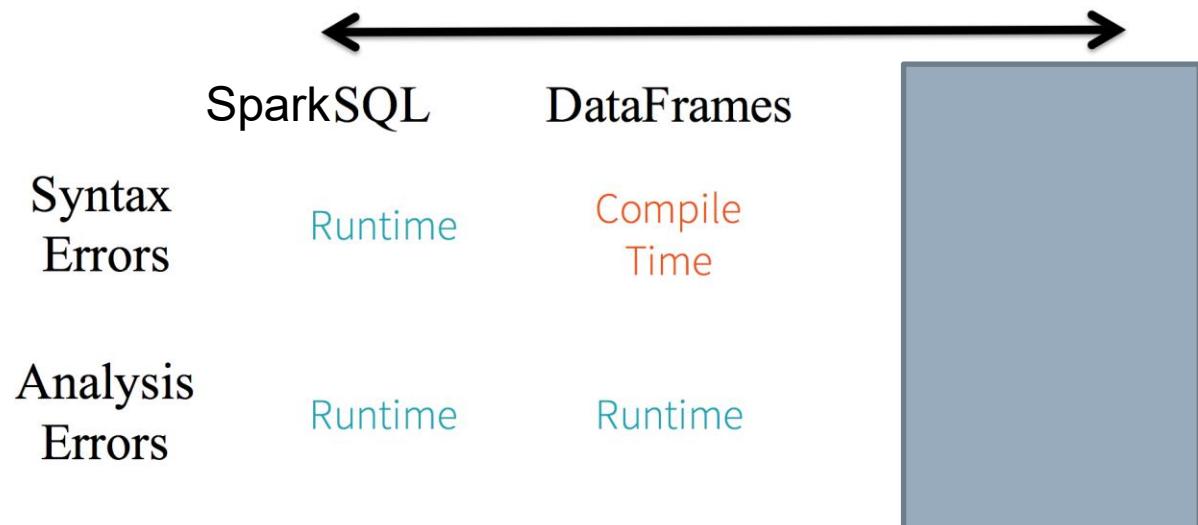


RDD vs. Datasets/DataFrames

- ▷ **RDD:** Immutable distributed collection
 - Transformations, actions, partitions
 - Lower-level abstraction. Better control. More coding required
- ▷ **DataFrame:** Immutable distributed collection (like RDD)
 - Data is organized into *named columns*
 - Imposes a structure, easier abstraction
 - Makes processing large data sets easier
 - Domain specific language API
- ▷ **Dataset:** Strongly-typed flavor of DataFrame
 - Only Java or Scala

RDD vs. Datasets/Data Frames

- ▷ Static-typing and runtime type-safety
 - Dataset easier to debug due to strong typing
 - RDD on Java/Scala is strongly typed
- ▷ High-level abstraction due to typing
- ▷ High-level abstraction due to columnar
 - RDD is row-based, Dataset/DF is both row & column based



RDD vs. Datasets/Data Frames

- ▶ Performance
 - Dataset and DataFrame use SparkSQL Catalyst optimizer...speed and space efficiency
 - Dataset efficient fast for *batch* execution. Specialized Tungsten serialization/deserialization and compact bytecode
 - DataFrame is untyped and faster for *interactive*
 - RDD gives fine-grained control
- ▶ You can move across them
 - `dataset.rdd.take(10)`

Declarative Programming

- ▷ Tell **what** to do, not *how* to do it
 - System determines the best plan
 - E.g., SQL and RDBMS query planning
 - **SELECT Part, Items FROM Widget WHERE Part='Bolts'**
- ▷ Uses schema for the data/table to plan the execution

RDD vs. DataFrames

```
# Create an RDD of tuples (name, age)
dataRDD = sc.parallelize([("Brooke", 20), ("Denny", 31), ("Jules", 30),
    ("TD", 35), ("Brooke", 25)])
# Use map and reduceByKey transformations with their lambda
# expressions to aggregate and then compute average

agesRDD = (dataRDD
    .map(lambda x: (x[0], (x[1], 1)))
    .reduceByKey(lambda x, y: (x[0] + y[0], x[1] + y[1])))
    .map(lambda x: (x[0], x[1][0]/x[1][1])))
```

```
# Create a DataFrame using SparkSession
spark = (SparkSession
    .builder
    .appName("AuthorsAges")
    .getOrCreate())
# Create a DataFrame
data_df = spark.createDataFrame([("Brooke", 20), ("Denny", 31), ("Jules", 30),
    ("TD", 35), ("Brooke", 25)], ["name", "age"])
# Group the same names together, aggregate their ages, and compute an average
avg_df = data_df.groupBy("name").agg(avg("age"))
# Show the results of the final execution
avg_df.show()
```

+-----+	name avg(age)	+-----+
Brooke 22.5		
Jules 30.0		
TD 35.0		
Denny 31.0		

DataFrames

- ▷ Distributed in-memory **tables** with named **columns** and **schemas**
 - Define *schema* explicitly by user
 - But **immutable** ... Spark keeps a lineage of all transformations
 - Add or change column names/types, creates new DataFrames while the previous versions are preserved

Table 3-1. The table-like format of a DataFrame

<i>Id</i> (Int)	<i>First</i> (String)	<i>Last</i> (String)	<i>Url</i> (String)	<i>Published</i> (Date)	<i>Hits</i> (Int)	<i>Campaigns</i> (List[Strings])
1	Jules	Damji	https://tinyurl.1	1/4/2016	4535	[twitter, LinkedIn]
2	Brooke	Wenig	https://tinyurl.2	5/5/2018	8908	[twitter, LinkedIn]
3	Denny	Lee	https://tinyurl.3	6/7/2019	7659	[web, twitter, FB, LinkedIn]
4	Tathagata	Das	https://tinyurl.4	5/12/2018	10568	[twitter, FB]

DataFrames: Define Schema using DDL

```
# Define schema for our data using DDL
schema = "`Id` INT, `First` STRING, `Last` STRING, `Url` STRING,
`Published` STRING, `Hits` INT, `Campaigns` ARRAY<STRING>

# Create our static data
data = [[1, "Jules", "Damji", "https://tinyurl.1", "1/4/2016", 4535, ["twitter",
"LinkedIn"]],
[2, "Brooke", "Wenig", "https://tinyurl.2", "5/5/2018", 8908, ["twitter",
"LinkedIn"]],
[3, "Denny", "Lee", "https://tinyurl.3", "6/7/2019", 7659, ["web",
"twitter", "FB", "LinkedIn"]],
[4, "Tathagata", "Das", "https://tinyurl.4", "5/12/2018", 10568,
["twitter", "FB"]],
[5, "Matei", "Zaharia", "https://tinyurl.5", "5/14/2014", 40578, ["web",
"twitter", "FB", "LinkedIn"]],
[6, "Reynold", "Xin", "https://tinyurl.6", "3/2/2015", 25568,
["twitter", "LinkedIn"]]
]

# Create a DataFrame using the schema defined above
blogs_df = spark.createDataFrame(data, schema)
# Show the DataFrame; it should reflect our table above
blogs_df.show()
# Print the schema used by Spark to process the DataFrame
print(blogs_df.printSchema())
```

Id	First	Last	Url	Published	Hits	Campaigns
1	Jules	Damji	https://tinyurl.1	1/4/2016	4535	[twitter,...]
2	Brooke	Wenig	https://tinyurl.2	5/5/2018	8908	[twitter,...]
3	Denny	Lee	https://tinyurl.3	6/7/2019	7659	[web, twitter,...]
4	Tathagata	Das	https://tinyurl.4	5/12/2018	10568	[twitter, FB]
5	Matei	Zaharia	https://tinyurl.5	5/14/2014	40578	[web, twitter,...]
6	Reynold	Xin	https://tinyurl.6	3/2/2015	25568	[twitter,...]

Projections & Filters

```
# In Python
```

```
few_fire_df = (fire_df
    .select("IncidentNumber", "AvailableDtTm", "CallType")
    .where(col("CallType") != "Medical Incident"))
few_fire_df.show(5, truncate=False)
```

IncidentNumber	AvailableDtTm	CallType
2003235	01/11/2002 01:47:00 AM	Structure Fire
2003235	01/11/2002 01:51:54 AM	Structure Fire
2003235	01/11/2002 01:47:00 AM	Structure Fire

```
# In Python, filter for only distinct non-null CallTypes from all the rows
```

```
(fire_df
    .select("CallType")
    .where(col("CallType").isNotNull())
    .distinct()
    .show(10, False))
```

CallType
Elevator / Escalator Rescue
Marine Fire
Aircraft Emergency
Confined Space / Structure Collapse

Aggregation & Join

```
# In Python
(fire_ts_df
    .select("CallType")
    .where(col("CallType").isNotNull())
    .groupBy("CallType")
    .count()
    .orderBy("count", ascending=False)
    .show(n=10, truncate=False))
```

Here, count() is the aggregator for groupBy() & not the action df.count()

CallType	count
Medical Incident	2843475
Structure Fire	578998
Alarms	483518
Traffic Collision	175507
Citizen Assist / Service Call	65360

```
# In Python
# Join departure delays data (foo) with airport info
foo.join(
    airports,
    airports.IATA == foo.origin
).select("City", "State", "date", "delay", "distance", "destination").show()
```

City	State	Country	IATA
Seattle	WA	USA	SEA
New York	NY	USA	JFK
Bangalore	KA	India	BLR

Origin	Dest-ination	Date	Delay	Distance	Airline
SEA	SFO	01010710	31	590	AA
SEA	SFO	01010955	104	590	UA
SEA	SFO	01010730	5	590	AA
LAX	SFO	01010600	15	400	DL

}

}

City	State	date	delay	distance	destination
Seattle	WA	01010710	31	590	SFO
Seattle	WA	01010955	104	590	SFO
Seattle	WA	01010730	5	590	SFO

Lazy Evaluation in DF

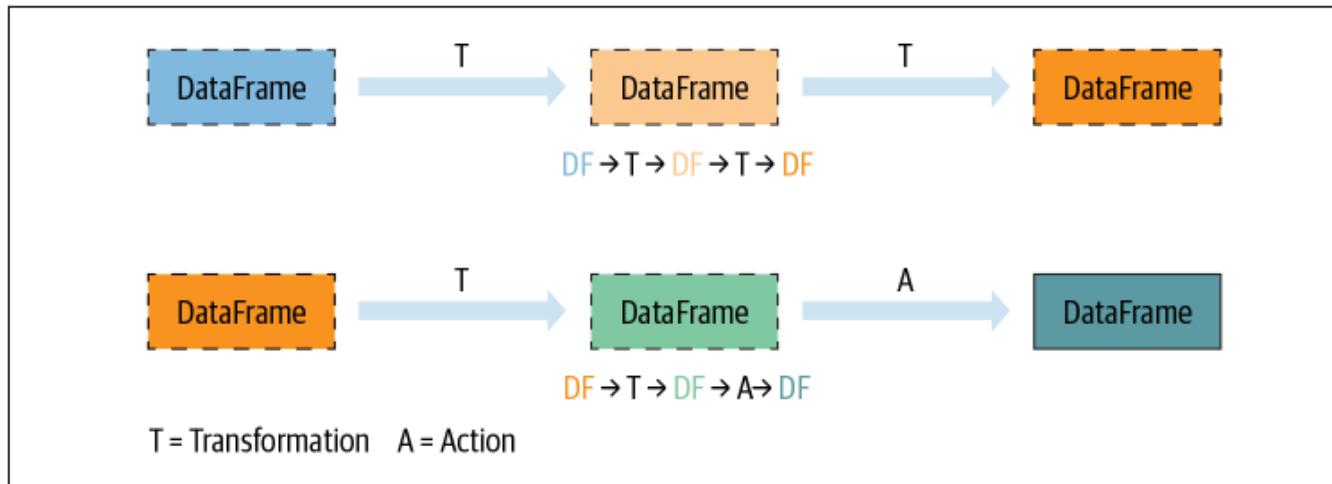


Figure 2-6. Lazy transformations and eager actions

Transformations	Actions
<code>orderBy()</code>	<code>show()</code>
<code>groupBy()</code>	<code>take()</code>
<code>filter()</code>	<code>count()</code>
<code>select()</code>	<code>collect()</code>
<code>join()</code>	<code>save()</code>

Immutability vs Modifications in DF

- ▷ DataFrames themselves are immutable
 - Backed by RDDs
- ▷ You can modify them to create new, different DataFrames
 - Add Columns: Original *foo* DataFrame plus the additional *status* derived column

```
# In Python
from pyspark.sql.functions import expr
foo2 = (foo.withColumn(
    "status",
    expr("CASE WHEN delay <= 10 THEN 'On-time' ELSE 'Delayed' END")
))


- Drop Cols: foo3 = foo2.drop("delay")
- Rename Cols: foo4 = foo3.withColumnRenamed("status", "flight_status")

```

- ▷ **ACID does not apply since immutable**

Spark SQL

- ▷ ANSI SQL:2003-compatible queries on structured data with a schema
- ▷ Permits abstraction to DataFrames/Datasets
- ▷ Connects to Apache Hive, JSON, CSV, Parquet,
- ▷ JDBC/ODBC and SQL Shell
- ▷ Optimized query plans

```
# In Python
count_mnm_df = (mnm_df
    .select("State", "Color", "Count")
    .groupBy("State", "Color")
    .agg(count("Count"))
    .alias("Total"))
    .orderBy("Total", ascending=False))
```

-- In SQL

```
SELECT State, Color, Count, sum(Count) AS Total
FROM MNM_TABLE_NAME
GROUP BY State, Color, Count
ORDER BY Total DESC
```

Count

Joins

```
# In Python
# Join departure delays data (foo) with airport info

foo.join(
    airports,
    airports.IATA == foo.origin
).select("City", "State", "date", "delay", "distance", "destination").show()

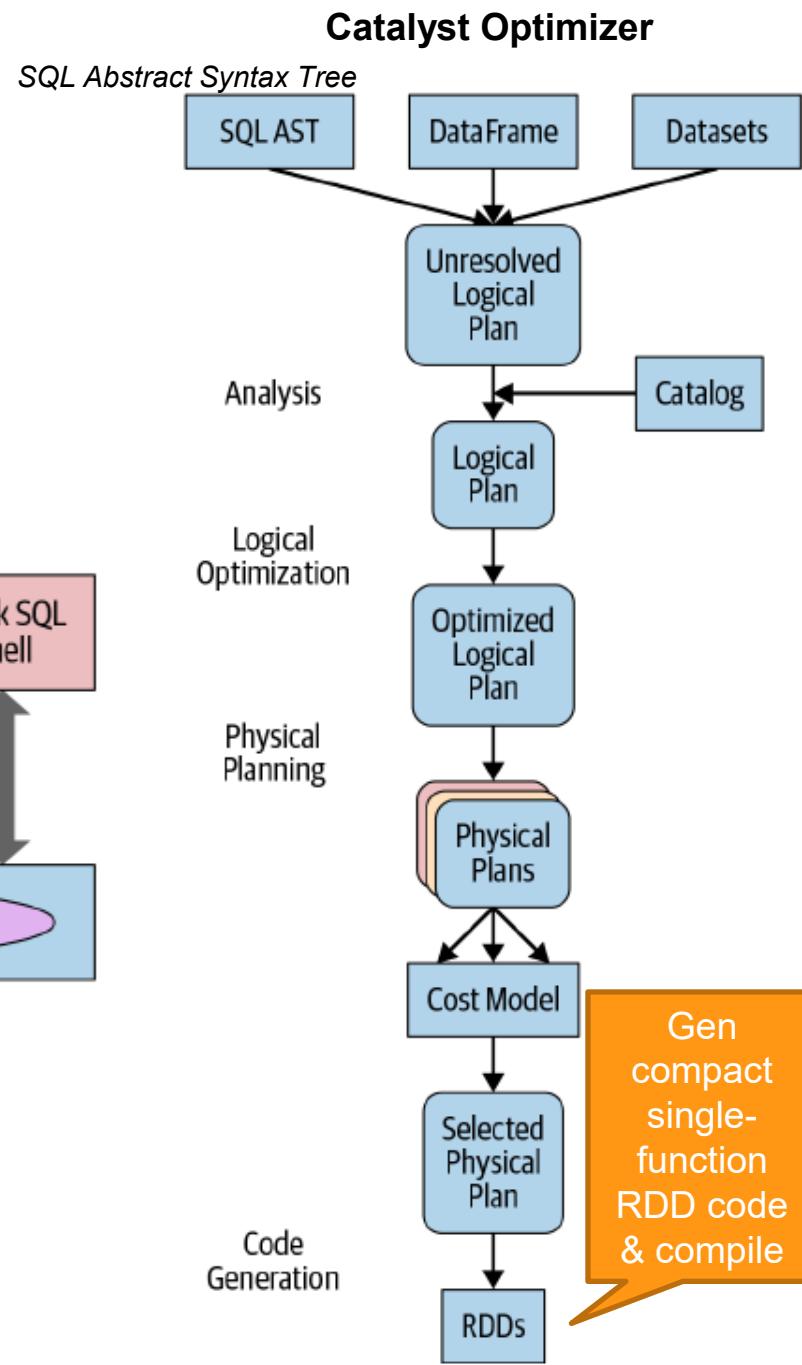
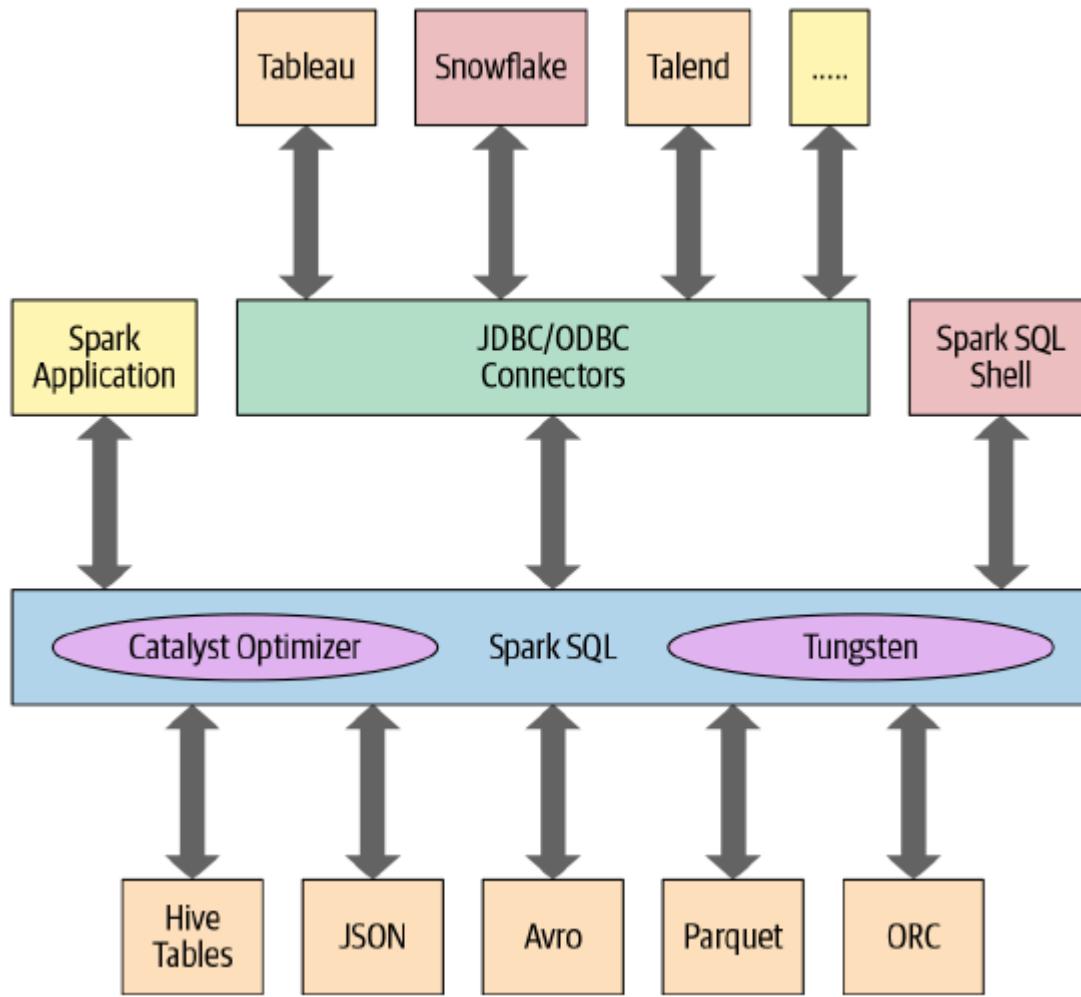
-- In SQL
spark.sql("""
SELECT a.City, a.State, f.date, f.delay, f.distance, f.destination
FROM foo f
JOIN airports_na a
    ON a.IATA = f.origin
""").show()
```

City	State	date	delay	distance	destination
Seattle	WA	01010710	31	590	SF0
Seattle	WA	01010955	104	590	SF0
Seattle	WA	01010730	5	590	SF0

Execution Model

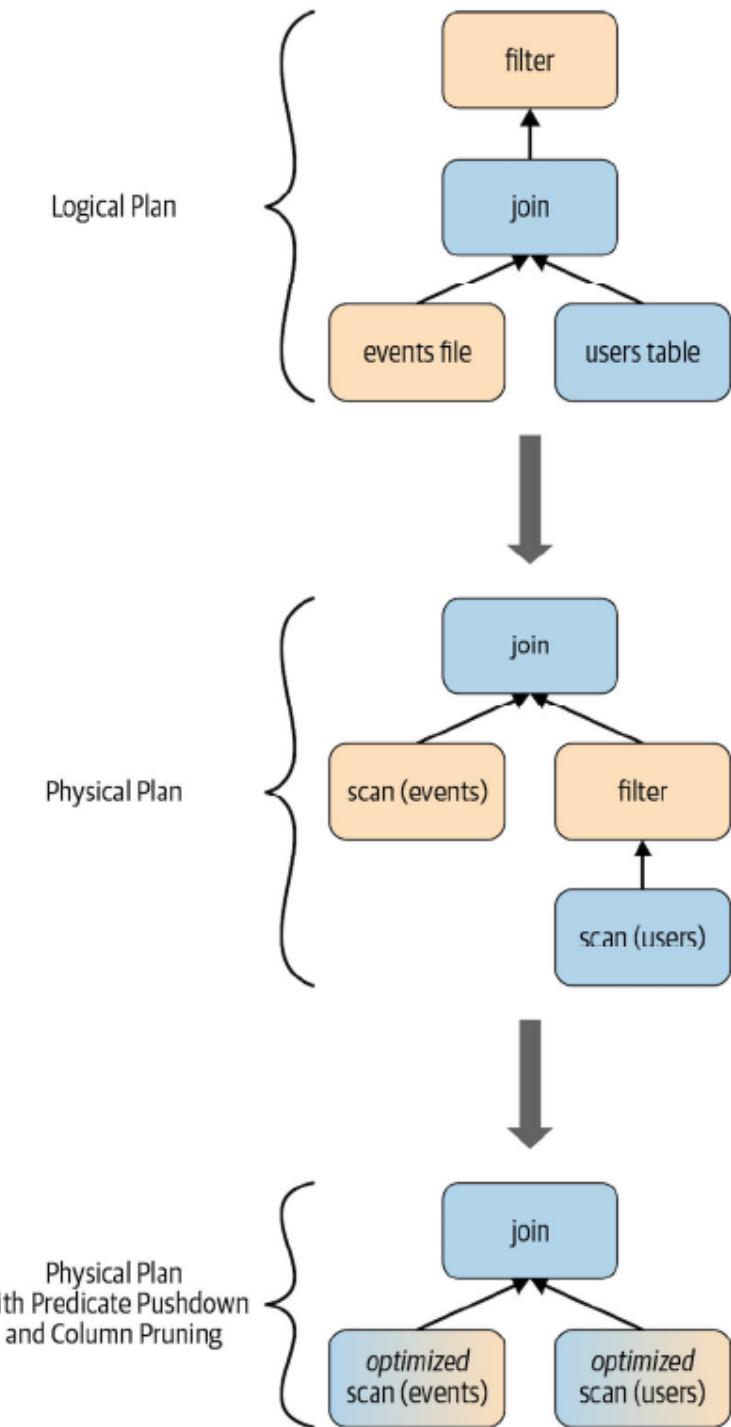
- ▶ Computation expressed in high-level DataFrame or SQL APIs is decomposed into low-level optimized and generated RDD operations
 - Then converted into Scala bytecode for the executors' JVMs
 - Generated RDD operation code is not accessible to users
 - RDD ops are NOT the same as the user-facing RDD APIs

SQL Stack & Query Plan

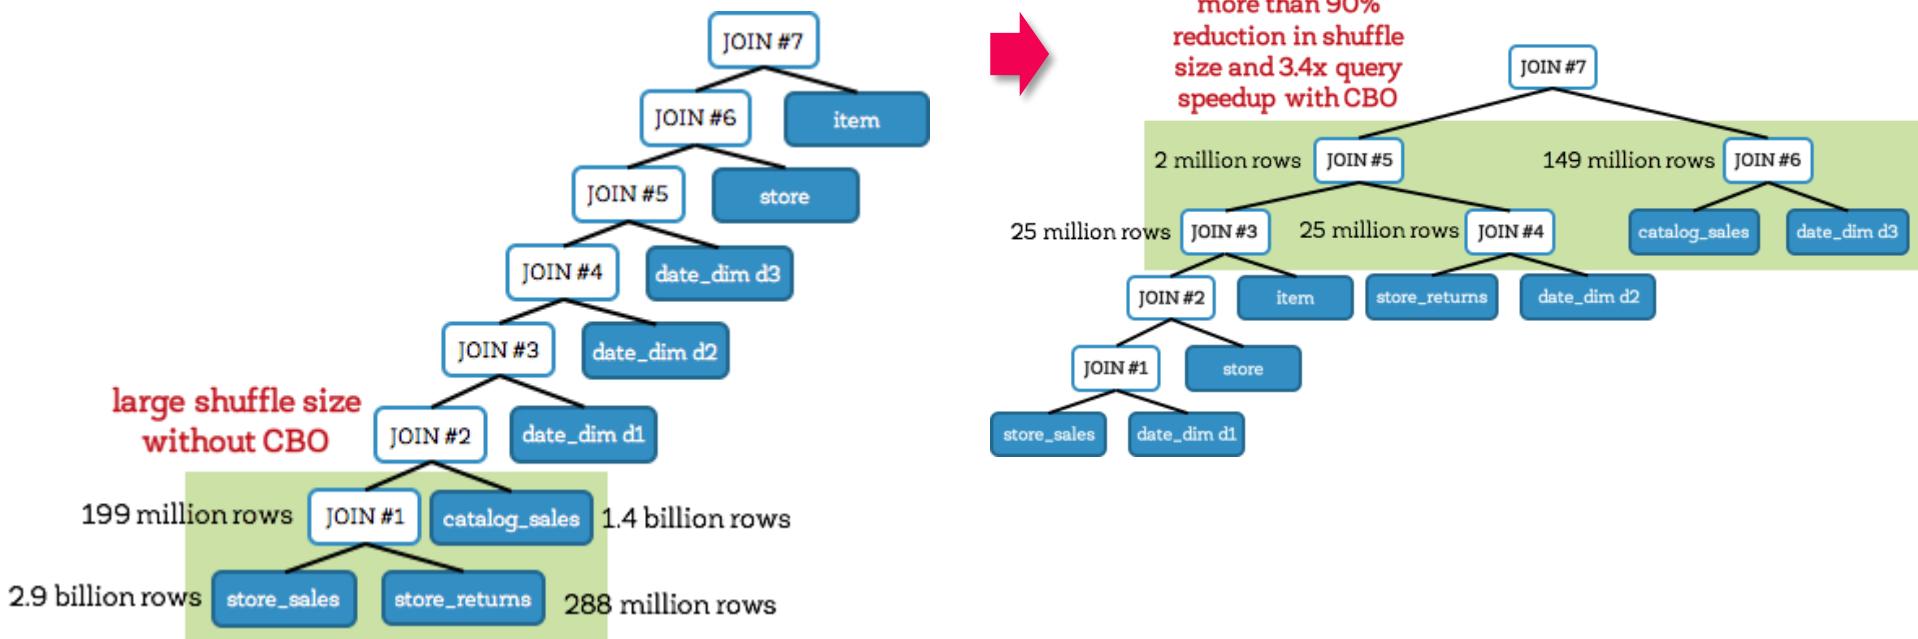


Query Planning

```
// Join two DataFrames
val joinedDF = users
  .join(events, users("id") === events("uid"))
  .filter(events("date") > "2015-01-01")
    users
```



Query Planning: Cost Based Optimizer



Catalyst Optimizer for Spark DF/SQL Execution

Spark SQL: Relational Data Processing in Spark, Michael Armbrust, et al., SIGMOD 2015.

(Advanced Topic)

Catalyst Optimizer

- ▷ DataFrames keep track of their schema
- ▷ DataFrames support various relational operations
 - DataFrame represents a *logical plan* to compute a dataset
 - No execution occurs (materialization of dataset) until an “action” output operation is called, such as save
- ▷ Enables rich optimization across all operations that were used to build the DataFrame
- ▷ User operations captured using abstract syntax tree rather than opaque Python/Scala functions

```
ctx = new HiveContext()
users = ctx.table("users")
young = users.where(users("age") < 21)
println(young.count())
```

Expression parsed by
Spark DF

DataFrame Domain Specific Language (DSL)

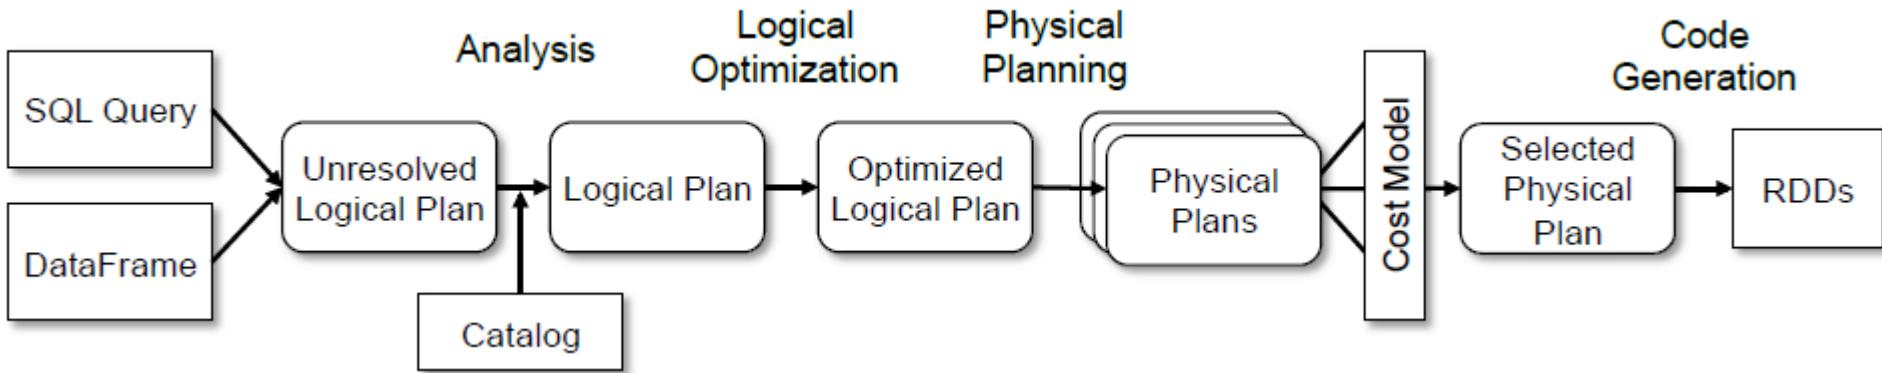
- ▷ **Relational Operators:** Projection (select), filter (where), join & aggregations (groupBy)
- ▷ Operators all take **expression** objects

```
employees
  .join(dept, employees("deptId") === dept("id"))
  .where(employees("gender") === "female")
  .groupBy(dept("id"), dept("name"))
  .agg(count("name"))
```

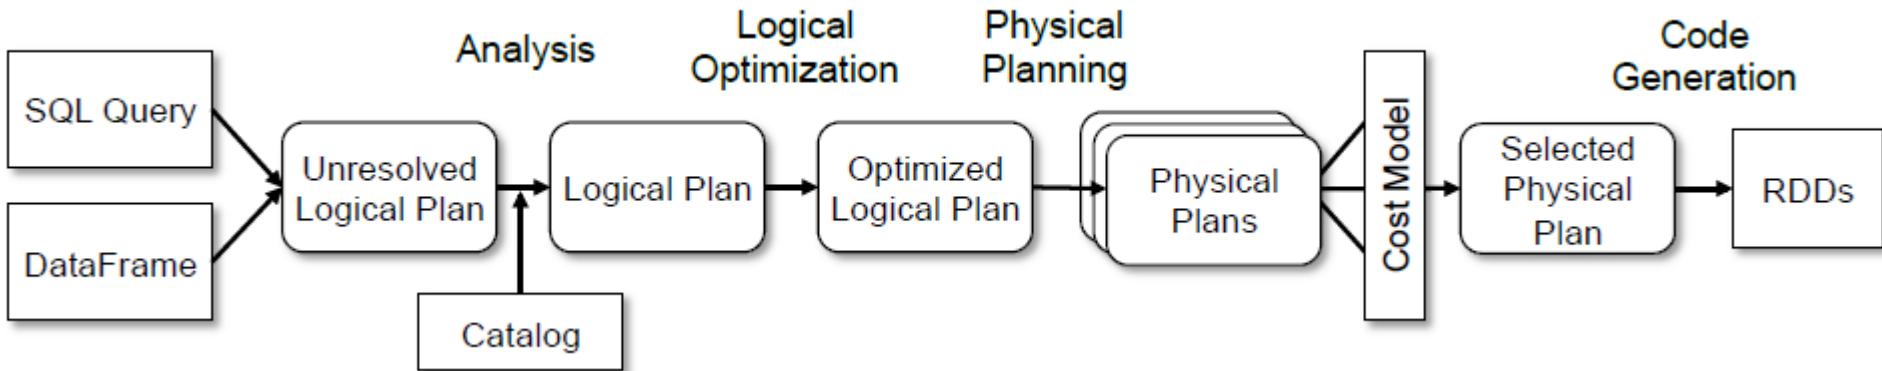
- ▷ Similar to SQL, but lets uses to split into logical steps, separate DFs, control flows, UDFs
 - Schema evaluation is eager to catch type errors
 - Execution is lazy when an action is seen

Catalyst Query Optimizer

- ▶ Modular and easy to add new optimization techniques and features to Spark SQL
- ▶ Enable external developers to extend the optimizer
- ▶ Cost-based Optimizer
 - Represent queries as **trees**
 - Applying **rules** to manipulate them
 - Generate different **plans** based on manipulation
 - Estimate the **execution cost** of each plan
 - Select the cheapest plan for actual execution

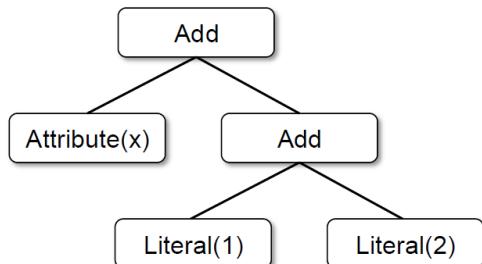


- ▷ Analyzing a logical plan to resolve references
 - AST from SQL or DF operations
 - Use *catalog* and inferencing to resolve attribute types, e.g., “SELECT (col+1) FROM sales”
 - Schema inference for semi-structured data



▷ Logical plan optimization

- Rule-based optimizations: Constant folding, Predicate pushdown, Projection pruning, Null propagation, Boolean expression simplification
- Trees and Rules for pattern matching/replacement
- Recursively apply rules till tree is simplified/optimized



```

tree.transform {
  case Add(Literal(c1), Literal(c2)) => Literal(c1+c2)
  case Add(left, Literal(0)) => left
  case Add(Literal(0), right) => right
}
  
```

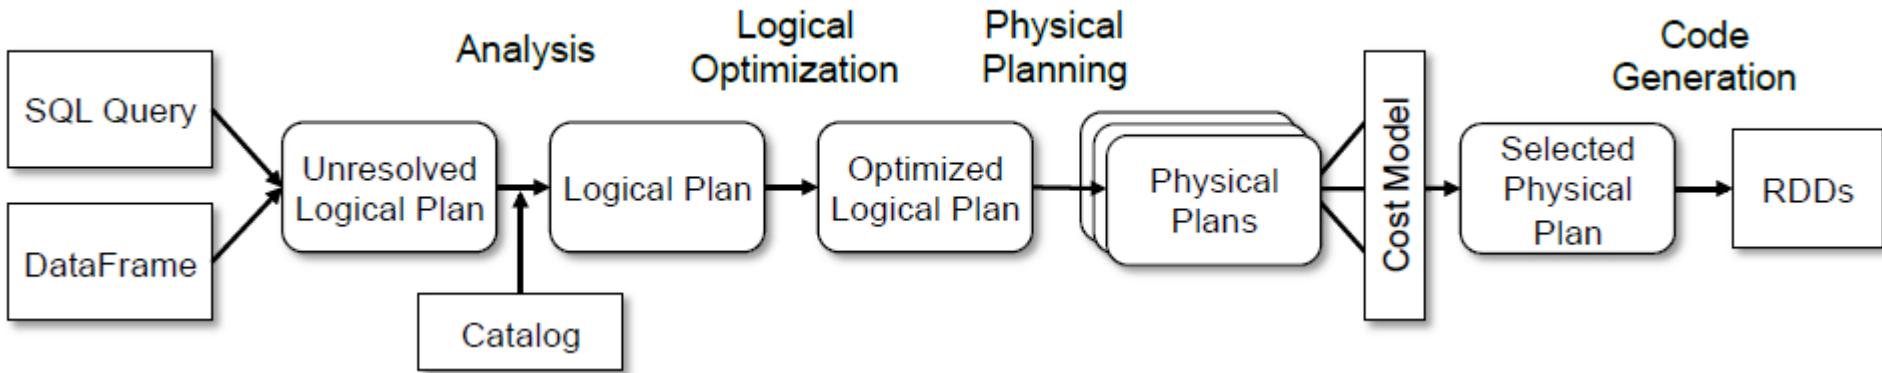
Float to Long for low precision math operations

```

object DecimalAggregates extends Rule[LogicalPlan] {
  /** Maximum number of decimal digits in a Long */
  val MAX_LONG_DIGITS = 18

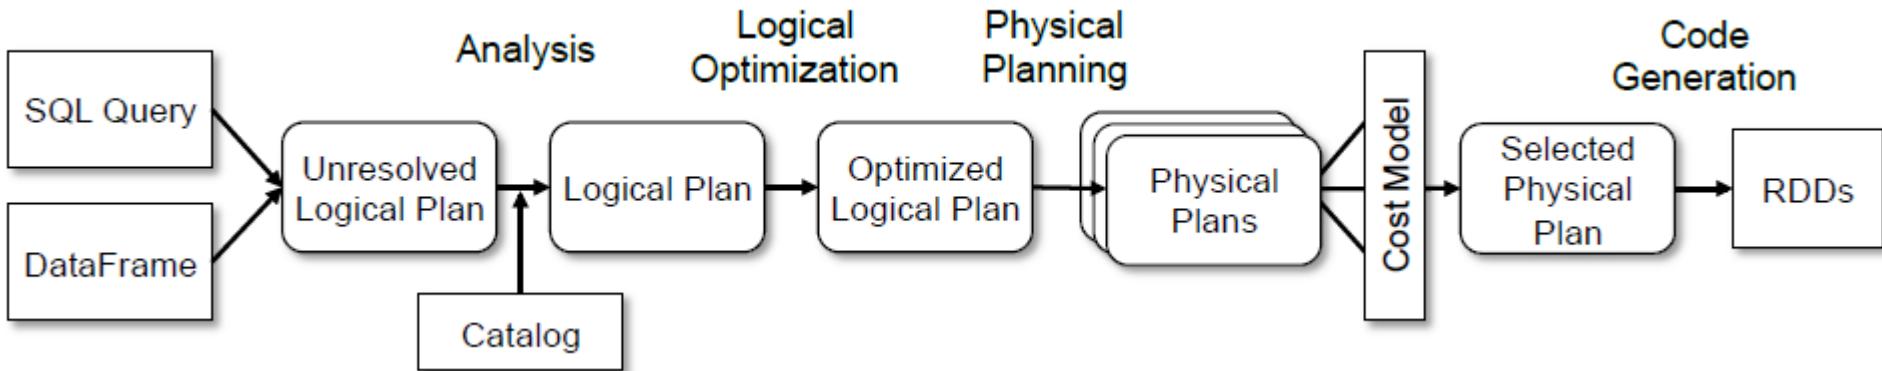
  def apply(plan: LogicalPlan): LogicalPlan = {
    plan.transformAllExpressions {
      case Sum(e @ DecimalType.Expression(prec, scale))
          if prec + 10 <= MAX_LONG_DIGITS =>
        MakeDecimal(Sum(LongValue(e)), prec + 10, scale)
    }
  }
}
  
```

Figure 2: Catalyst tree for the expression $x+(1+2)$.



▷ Physical planning

- Takes a logical plan and generates one or more physical plans
- Uses physical operators present in the Spark (RDD) execution engine
- Rule-based optimizations: Pipelining projections or filters into one Spark RDD **map** operation
- Cost based optimization: Estimates table sizes using in-memory cache, external file sizing, result of a subquery with a LIMIT



- ▷ Code generation to compile parts of the query to Java bytecode
 - Operates on in-memory datasets, so processing is *CPU-bound*
 - Code generation speeds up execution
 - Interpreting and evaluating each expression using if/else/switch is costly
 - AST Converted to Scala code, compiled into byte-code



“

- ▷ *Until here for Test #1 on 12 Feb on
Module 1 and 2*

“

End of Lecture 11

Design Patterns

HeartBeat

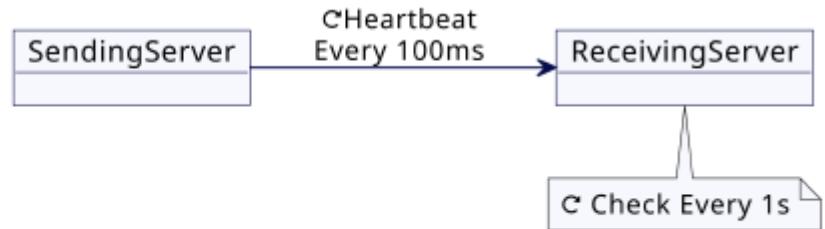


Figure 7.1 *Heartbeat*

- ▷ Timely detection of server failures
- ▷ timeout interval > request interval > network round trip time between the servers

From Patterns of Distributed Systems, Unmesh Joshi, Martin Fowler, 2023

Design Patterns

A Consistent Core Can Manage the Membership of a Data Cluster

- ▶ HDFS Name Node, later Kafka

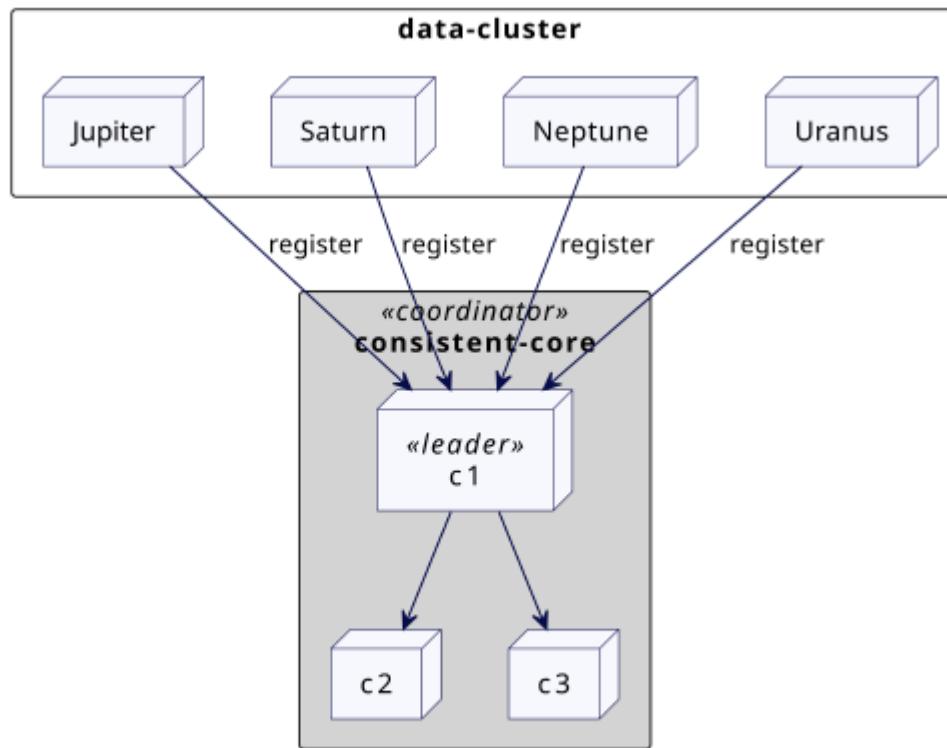


Figure 2.53 *Consistent Core tracks cluster membership.*

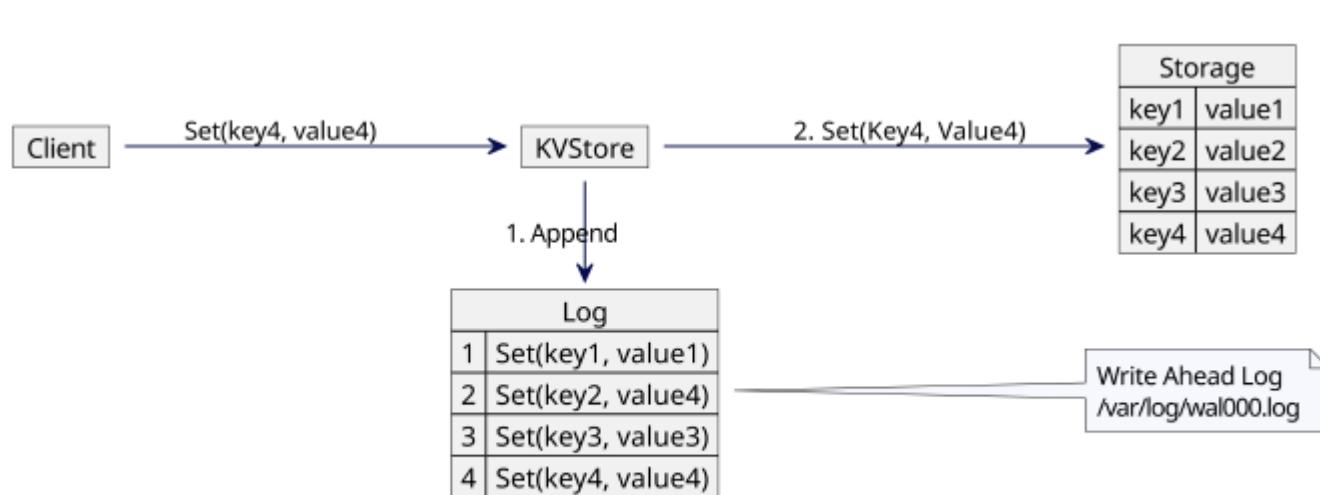
Design Patterns

Leasing

- ▷ A node can ask for a lease for a limited period of time, after which it expires.
- ▷ The node can renew the lease before it expires if it wants to extend the access.
- ▷ Implement the lease mechanism with *Consistent Core* to provide fault tolerance and consistency.
- ▷ Have a time-to-live value associated with the lease.
- ▷ Leader node tracks the lease timeouts, using its own monotonic clock...*why not wallclock?*
- ▷ *HDFS Name node and leases to Block leaders*

Design Patterns

Keeping Data Resilient on a Single Server



A variation used in Name Node of HDFS

Figure 3.1 Write-Ahead Log

Design Patterns

Replicating Client Requests

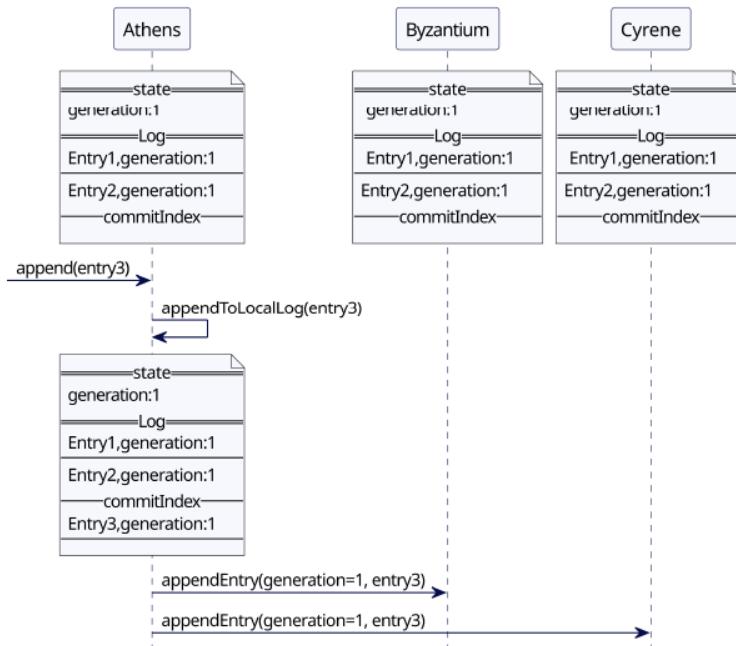


Figure 12.1 *Leader appends to its own log.*

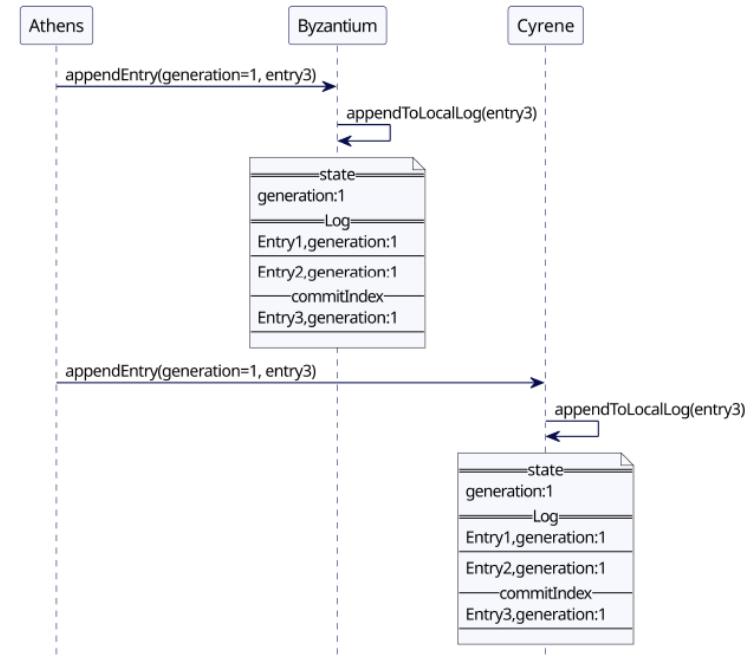


Figure 12.2 *Followers append to their logs.*

A Large Amount of Data Can Be Partitioned over Multiple Nodes

- ▶ Data should be evenly distributed across all the cluster nodes.
 - Fixed block size (HDFS) vs Variable block size (RDD)
- ▶ It should be possible to know which cluster node stores a particular data record, without making a request to all the nodes.
 - Costly Lookups (HDFS) vs Static Hashing (RDD shuffle)
- ▶ It should be quick and easy to move part of the data to the new nodes.
 - Logical (Dyanamo) vs Physical Partitioning (HDFS)

Design Patterns

Partitions Can Be Replicated for Resilience

- ▷ All Replication (HDFS Data Node) vs. Majority Quorum Replication
 - Leader election

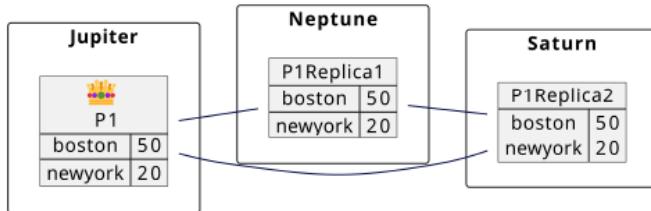


Figure 2.40 Partitions are replicated for resilience.

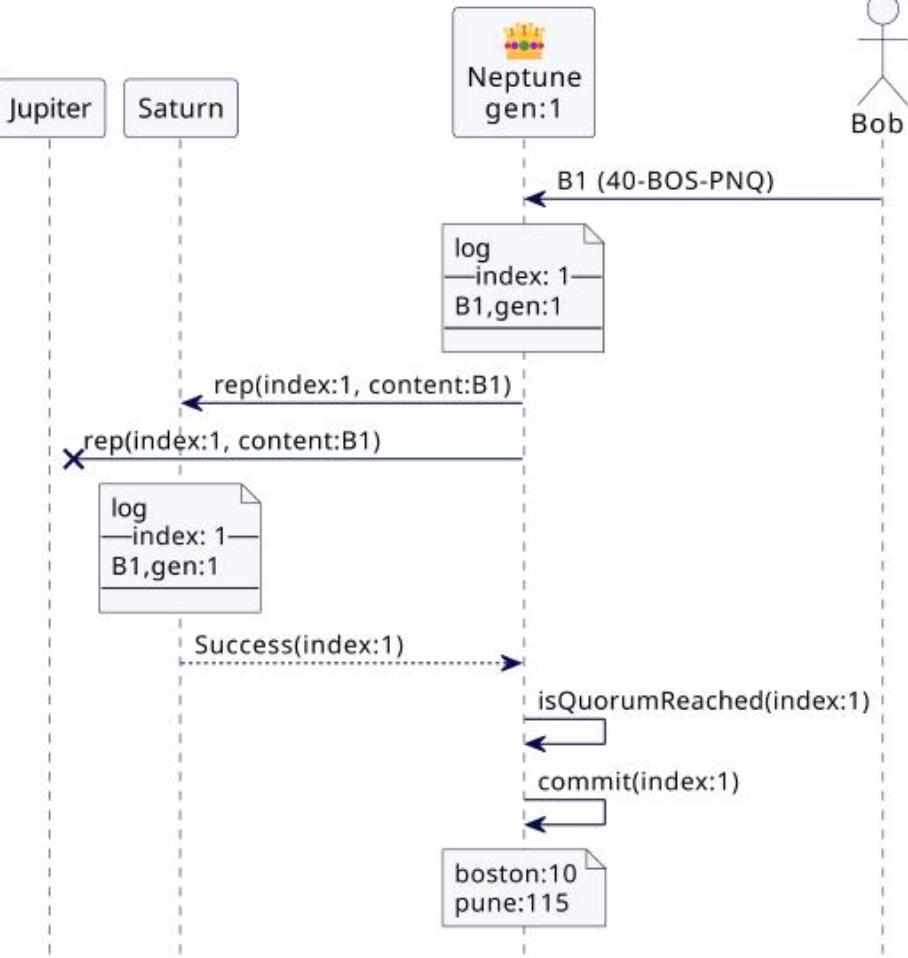


Figure 2.20 Log entries are committed once they are accepted by a Majority Quorum.

Design Patterns

Versioning

- ▷ HDFS uses block versioning
 - Prevent DataNodes from serving outdated or corrupted block replicas
- ▷ <BlockID, GenerationStamp> on NameNode
 - GenerationStamp (GS): Monotonically increasing version number assigned by NameNode upon Write or Append
 - Higher GS → newer version
- ▷ Any DataNode replica with a lower GS is stale
 - During the write, replicas must all have the same GS

Design Patterns

Request Pipeline

- ▷ Send next request without waiting for previous request to be acked.
- ▷ Reduces waiting time latency.
- ▷ Queue to buffer pending requests.
- ▷ *Data node write pipeline in HDFS...*

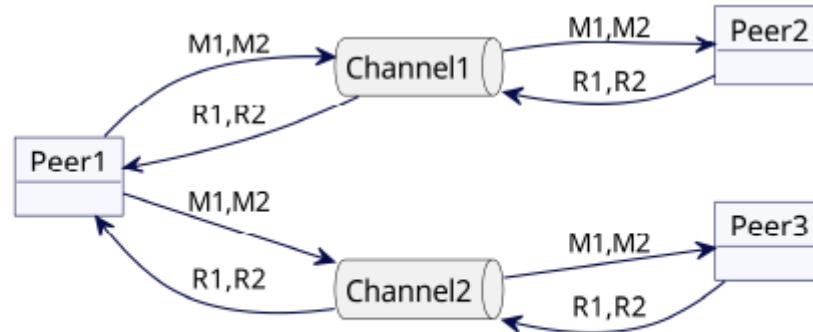


Figure 32.1 Request pipeline

Clash of the Titans: MapReduce vs. Spark for Large Scale Data Analytics

Juwei Shi, Yunjie Qiu, Umar Farooq Minhas, Limei Jiao Chen
Wang, Berthold Reinwald, and Fatma Ozcan
VLDB 2015

- Hadoop version 2.4.0
- Spark version 1.3.0

Architectural Diffs

- ▷ Evaluation of *batch* and *iterative* jobs
- ▷ **Shuffle stage**
 - MR: Shuffle+Sort, combiner/aggregation, external merge
 - Spark: Hash shuffle
- ▷ **Execution**
 - Task parallelism, Overlapping computation, Data Pipelining among stages
- ▷ **Caching**
 - Reuse of intermediate data. OS, HDFS, Tachyon, RDD

Workload

Table 1: Characteristics of Selected Workloads

		Word Count	Sort	K-Means (LR)	Page-Rank
Type	One Pass	✓	✓		
	Iterative			✓	✓
Shuffle Sel.	High		✓		
	Medium				✓
	Low	✓		✓	
Job/Iter. Sel.	High		✓		
	Medium				✓
	Low	✓		✓	

Ratio of Map output size to Job input size

Ratio of Reduce output size to Job input size

Workload

Table 2: Key Architectural Components of Interest

Incl. Combiner		Word Count	Sort	K-Means (LR)	Page-Rank
Merge sort for large shuffle data	Shuffle	Aggregation	✓		✓
		External sort	✓		
		Data transfer	✓		✓
Compute+ NW xfr	Execution	Task parallelism	✓	✓	✓
		Stage overlap	✓		
		Data pipelining			✓
Reuse w/o HDFS write	Caching	Input		✓	✓
		Intermediate data			✓

Setup

- ▷ 4 servers @ 32 CPU cores, 2.9GHz,
 - 9 disk drives at 7.2k RPM with 1 TB each
 - Hard disks deliver an aggregate bandwidth of about 125 GB/sec for reads and 45 GB/sec for writes
 - 190GB RAM
- ▷ 1 Gbps Ethernet switch
- ▷ We use Hadoop version 2.4.0 to run MapReduce on YARN
 - ▷ 32 containers per node
- ▷ We use Spark version 1.3.0 running in the standalone mode on HDFS 2.4.0
 - ▷ 8 Spark workers per node with 4 threads each

Summary: Diffs

- ▷ Performance differences due to components architectures in the two frameworks.
- ▷ Spark is about 2.5x, 5x, and 5x faster than MapReduce, for Word Count, k-means, and PageRank
 - + efficiency of the **hash-based aggregation** component for **combine** (Map-side reduction), 40% improvement
 - + **reduced CPU and disk overheads** due to RDD caching in Spark (PR, k-Means), 90%
 - + Data pipelining, avoid materialization
 - + Task loading , **context switch** is 10x faster

Summary: Diffs

- ▷ MapReduce is 2x faster than Spark for Sort workload
 - + MapReduce is **more efficient for shuffling data** than Spark, overlap Shuffle with Map, hide network overhead
 - Map stage in Spark is slower with more Reducers due to more open files
 - Increasing JVM for Spark has GC overheads

Summary: Common

- ▷ For one pass jobs, **Map is CPU bound, Reduce is network bound**
 - So **disk I/O not a bottleneck** (NW is), so spills often do not have a lot of penalty
- ▷ **Input parsing** is often an overhead
 - RDD caching helps. OS/HDFS caching does not.
- ▷ **GC overhead** bottleneck if heap size per task drops to 64MB with 128MB split
- ▷ Disk caching is bottleneck for RDD if CPU and disk I/O capacities unbalanced